

Assignment 1a: Machine Learning Methods

\yuqing chen \ a1841612

How can we use data analytics to predict energy usage for this house?



Follow detail instructions in the assignment specification and use these heading for implementation and discussion

1. Read the dataset

```
In [ ]: import pandas as pd
import numpy as np

# load the data from the dataset
file_path = 'energydata_complete.csv'
df = pd.read_csv(file_path)
```

```
/var/folders/n3/ddr1nygd4b11p0zv8kwtqx3w0000gn/T/ipykernel_94042/148841221
3.py:1: DeprecationWarning:
Pyarrow will become a required dependency of pandas in the next major rele
ase of pandas (pandas 3.0),
(to allow more performant data types, such as the Arrow string type, and b
etter interoperability with other libraries)
but was not found to be installed on your system.
If this would cause problems for you,
please provide us feedback at https://github.com/pandas-dev/pandas/issues/
54466
```

```
import pandas as pd
```

```
In [ ]: # check if the data is loaded successfully
print(df.head())
```

		date	Appliances	lights	T1	RH_1	T2	RH_2	
\									
0	11/01/2016 17:00		60	30	19.89	47.596667	19.2	44.790000	
1	11/01/2016 17:10		60	30	19.89	46.693333	19.2	44.722500	
2	11/01/2016 17:20		50	30	19.89	46.300000	19.2	44.626667	
3	11/01/2016 17:30		50	40	19.89	46.066667	19.2	44.590000	
4	11/01/2016 17:40		60	40	19.89	46.333333	19.2	44.530000	
		T3	RH_3	T4	...	T9	RH_9	T_out	Press_mm_hg
\									
0	19.79	44.730000	19.000000	...	17.033333	45.53	6.60		733.5
1	19.79	44.790000	19.000000	...	17.066667	45.56	6.48		733.6
2	19.79	44.933333	18.926667	...	17.000000	45.50	6.37		733.7
3	19.79	45.000000	18.890000	...	17.000000	45.40	6.25		733.8
4	19.79	45.000000	18.890000	...	17.000000	45.40	6.13		733.9
		RH_out	Windspeed	Visibility	Tdewpoint	rv1	rv2		
0	92.0	7.000000	63.000000		5.3	13.275433	13.275433		
1	92.0	6.666667	59.166667		5.2	18.606195	18.606195		
2	92.0	6.333333	55.333333		5.1	28.642668	28.642668		
3	92.0	6.000000	51.500000		5.0	45.410390	45.410390		
4	92.0	5.666667	47.666667		4.9	10.084097	10.084097		

[5 rows x 29 columns]

Now we can see the data is successfully loaded. Next we will analyse the data

2. Analyse and visualise the data

Check the data information and the types

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19735 entries, 0 to 19734
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                   19735 non-null  object
1   Appliances              19735 non-null  int64
2   lights                  19735 non-null  int64
3   T1                      19735 non-null  float64
4   RH_1                    19735 non-null  float64
5   T2                      19735 non-null  float64
6   RH_2                    19735 non-null  float64
7   T3                      19735 non-null  float64
8   RH_3                    19735 non-null  float64
9   T4                      19735 non-null  float64
10  RH_4                    19735 non-null  float64
11  T5                      19735 non-null  float64
12  RH_5                    19735 non-null  float64
13  T6                      19735 non-null  float64
14  RH_6                    19735 non-null  float64
15  T7                      19735 non-null  float64
16  RH_7                    19735 non-null  float64
17  T8                      19735 non-null  float64
18  RH_8                    19735 non-null  float64
19  T9                      19735 non-null  float64
20  RH_9                    19735 non-null  float64
21  T_out                   19735 non-null  float64
22  Press_mm_hg            19735 non-null  float64
23  RH_out                  19735 non-null  float64
24  Windspeed               19735 non-null  float64
25  Visibility              19735 non-null  float64
26  Tdewpoint               19735 non-null  float64
27  rv1                     19735 non-null  float64
28  rv2                     19735 non-null  float64
dtypes: float64(26), int64(2), object(1)
memory usage: 4.4+ MB
```

We can see the type of the date is an object which is different from others, so we need to do convert. According to research, dataset's date/time variables can generate additional functions: the number of seconds from midnight for each day (NSM), the week status (weekend or weekday) and the day of the week (Candanedo et al. 2017, P. 85). So I will convert the 'date' time variable to NSM and week status

```
In [ ]: import pandas as pd

# Specify the datetime format as '%d/%m/%Y %H:%M'
df['date'] = pd.to_datetime(df['date'], format='%d/%m/%Y %H:%M')

# Generate NSM (number of seconds per day from midnight)
df['NSM'] = (df['date'] - df['date'].dt.normalize()).dt.total_seconds()

# status of weekend and weekday (0 = weekday, 1 = weekend)
df['weekend'] = df['date'].dt.dayofweek // 5

# load the data and check if these two columns are correctly added
print(df.head(700))
```

	date	Appliances	lights	T1	RH_1	T2	\
0	2016-01-11 17:00:00	60	30	19.890000	47.596667	19.2	
1	2016-01-11 17:10:00	60	30	19.890000	46.693333	19.2	
2	2016-01-11 17:20:00	50	30	19.890000	46.300000	19.2	
3	2016-01-11 17:30:00	50	40	19.890000	46.066667	19.2	
4	2016-01-11 17:40:00	60	40	19.890000	46.333333	19.2	
..	
695	2016-01-16 12:50:00	90	0	22.075000	40.045000	21.3	
696	2016-01-16 13:00:00	90	0	22.000000	39.223333	21.1	
697	2016-01-16 13:10:00	100	0	22.000000	39.626667	21.1	
698	2016-01-16 13:20:00	80	0	22.000000	39.566667	21.1	
699	2016-01-16 13:30:00	90	0	21.963333	39.560000	21.1	

	RH_2	T3	RH_3	T4	...	T_out	Press_mm_hg	\
0	44.790000	19.790000	44.730000	19.000000	...	6.60	733.500000	
1	44.722500	19.790000	44.790000	19.000000	...	6.48	733.600000	
2	44.626667	19.790000	44.933333	18.926667	...	6.37	733.700000	
3	44.590000	19.790000	45.000000	18.890000	...	6.25	733.800000	
4	44.530000	19.790000	45.000000	18.890000	...	6.13	733.900000	
..	
695	38.923333	21.033333	41.723333	20.196667	...	5.70	763.733333	
696	37.950000	20.890000	40.990000	20.100000	...	5.90	763.700000	
697	38.156667	20.963333	40.790000	20.166667	...	5.88	763.700000	
698	38.363333	21.000000	40.826667	20.166667	...	5.87	763.700000	
699	38.500000	21.000000	40.900000	20.166667	...	5.85	763.700000	

	RH_out	Windspeed	Visibility	Tdewpoint	rv1	rv2	\
0	92.000000	7.000000	63.000000	5.300	13.275433	13.275433	
1	92.000000	6.666667	59.166667	5.200	18.606195	18.606195	
2	92.000000	6.333333	55.333333	5.100	28.642668	28.642668	
3	92.000000	6.000000	51.500000	5.000	45.410390	45.410390	
4	92.000000	5.666667	47.666667	4.900	10.084097	10.084097	
..	
695	69.500000	4.666667	26.333333	0.450	27.965514	27.965514	
696	66.000000	5.000000	27.000000	0.000	39.454724	39.454724	
697	68.333333	4.833333	29.166667	0.417	8.385763	8.385763	
698	70.666667	4.666667	31.333333	0.833	48.522587	48.522587	
699	73.000000	4.500000	33.500000	1.250	23.675155	23.675155	

	NSM	weekend
0	61200.0	0
1	61800.0	0
2	62400.0	0
3	63000.0	0
4	63600.0	0
..
695	46200.0	1
696	46800.0	1
697	47400.0	1
698	48000.0	1
699	48600.0	1

[700 rows x 31 columns]

After the converting, check the info of the dataset again.

```
In [ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19735 entries, 0 to 19734
Data columns (total 31 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   date                  19735 non-null  datetime64[ns]
 1   Appliances            19735 non-null  int64
 2   lights               19735 non-null  int64
 3   T1                   19735 non-null  float64
 4   RH_1                 19735 non-null  float64
 5   T2                   19735 non-null  float64
 6   RH_2                 19735 non-null  float64
 7   T3                   19735 non-null  float64
 8   RH_3                 19735 non-null  float64
 9   T4                   19735 non-null  float64
10  RH_4                 19735 non-null  float64
11  T5                   19735 non-null  float64
12  RH_5                 19735 non-null  float64
13  T6                   19735 non-null  float64
14  RH_6                 19735 non-null  float64
15  T7                   19735 non-null  float64
16  RH_7                 19735 non-null  float64
17  T8                   19735 non-null  float64
18  RH_8                 19735 non-null  float64
19  T9                   19735 non-null  float64
20  RH_9                 19735 non-null  float64
21  T_out                19735 non-null  float64
22  Press_mm_hg          19735 non-null  float64
23  RH_out               19735 non-null  float64
24  Windspeed            19735 non-null  float64
25  Visibility            19735 non-null  float64
26  Tdewpoint            19735 non-null  float64
27  rv1                  19735 non-null  float64
28  rv2                  19735 non-null  float64
29  NSM                  19735 non-null  float64
30  weekend                19735 non-null  int32
dtypes: datetime64[ns](1), float64(27), int32(1), int64(2)
memory usage: 4.6 MB

```

We can see two variables were successfully added

Now construct a heat map of the correlation between features

```

In [ ]: import seaborn as sns
import matplotlib.pyplot as plt

correlation_matrix = df.corr()

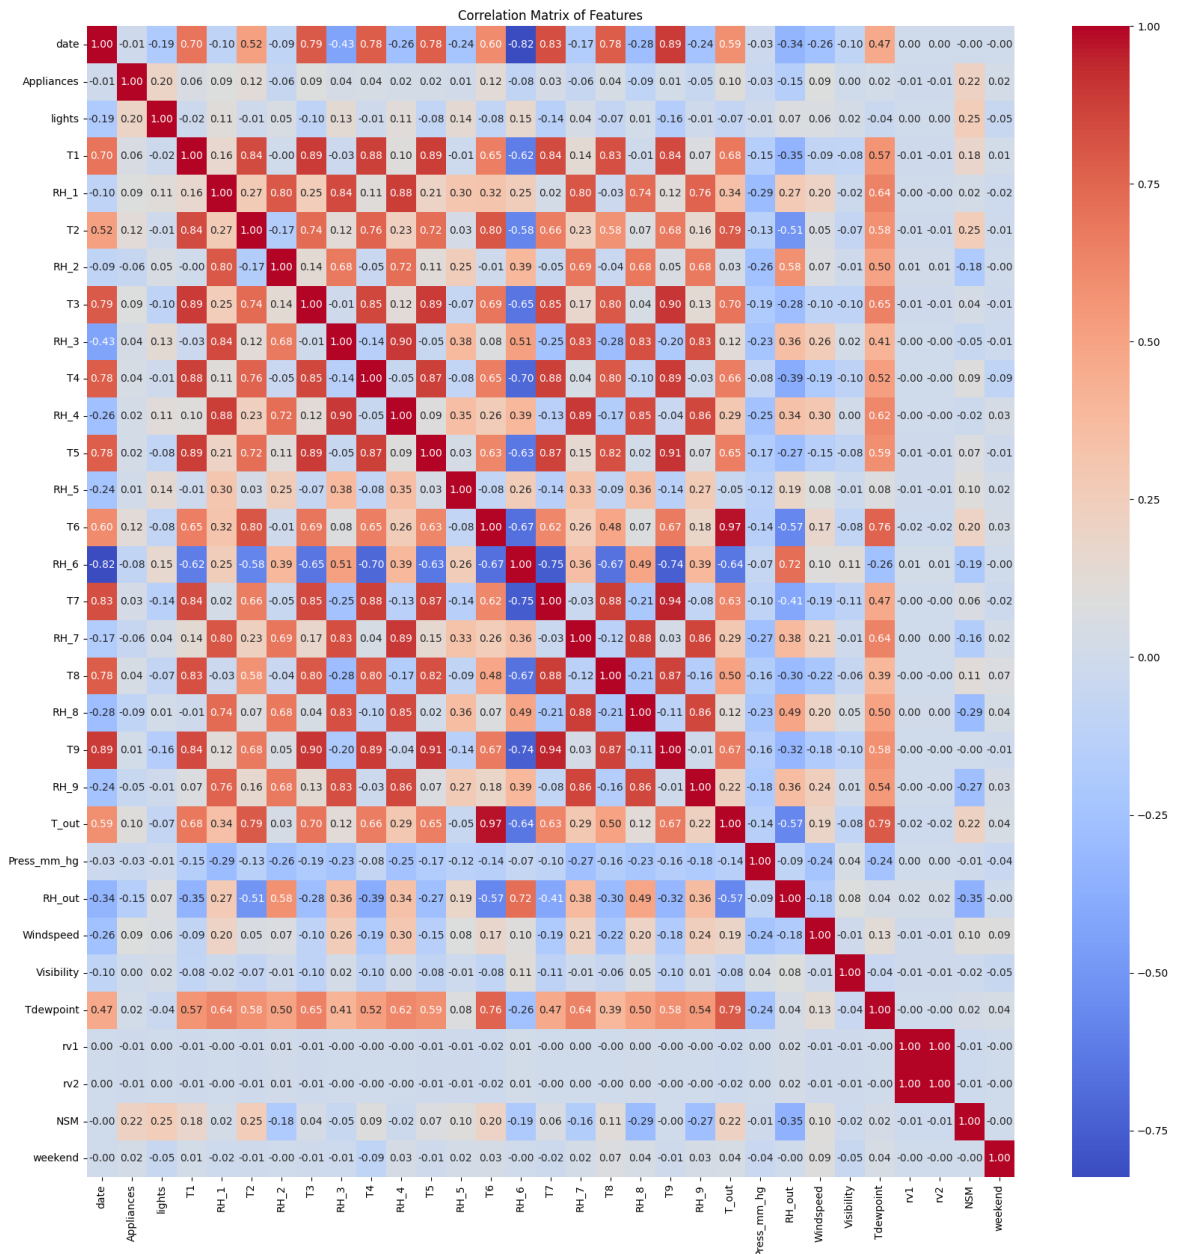
plt.figure(figsize=(20, 20))

sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")

# add the title
plt.title('Correlation Matrix of Features')

plt.show()

```



According to the heatmap, Appliance is highly correlated with light and NSM. Next, I will analyse the appliance variable

```
In [ ]: # check the appliances column information
df['Appliances'].describe()
```

```
Out[ ]: count      19735.000000
mean         97.694958
std         102.524891
min          10.000000
25%          50.000000
50%          60.000000
75%         100.000000
max         1080.000000
Name: Appliances, dtype: float64
```

To find outliers, I will visualise the appliances variables as line graphs, histograms and box plots


```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(20, 6))
plt.title('appliances consumption over Time')
plt.plot_date(df['date'], df['Appliances'], linestyle='solid', marker = 'o')
plt.gcf().autofmt_xdate()

plt.xticks(df['date'][:4000]) # Display a date every 4000 data points
plt.show()
```

/var/folders/n3/ddr1nygd4b11p0zv8kwtqx3w0000gn/T/ipykernel_94042/4165792070.py:6: UserWarning: marker is redundantly defined by the 'marker' keyword argument and the fmt string "o" (-> marker='o'). The keyword argument will take precedence.

```
plt.plot_date(df['date'], df['Appliances'], linestyle='solid', marker = 'None')
```



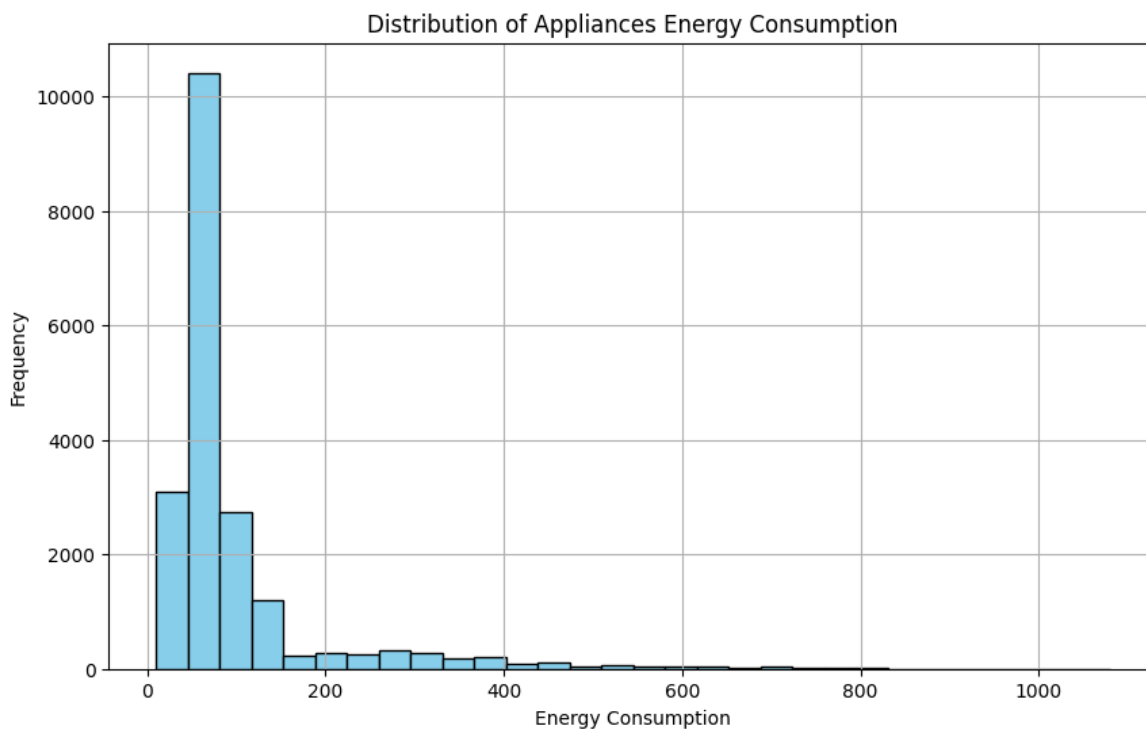
```
In [ ]: plt.figure(figsize=(10, 6))

plt.hist(df['Appliances'], bins=30, color='skyblue', edgecolor='black')

plt.title('Distribution of Appliances Energy Consumption')
plt.xlabel('Energy Consumption')
plt.ylabel('Frequency')

plt.grid(True)

plt.show()
```



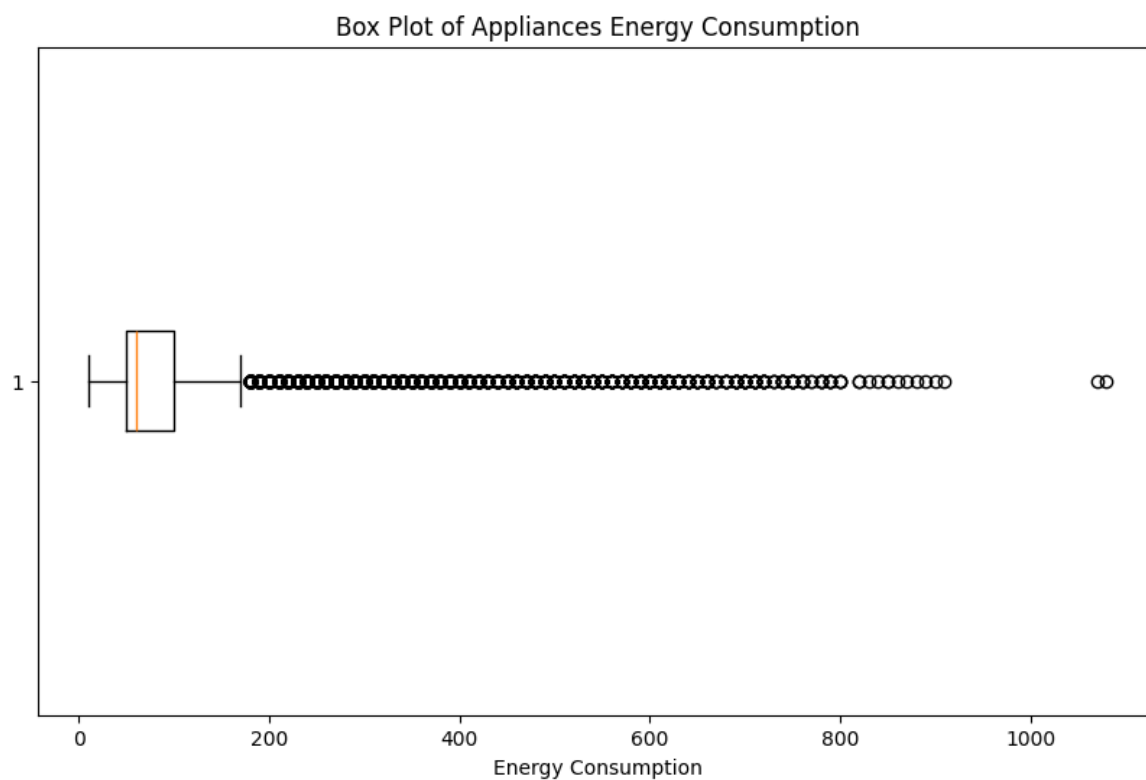
```
In [ ]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))

plt.boxplot(df['Appliances'], vert=False)

plt.title('Box Plot of Appliances Energy Consumption')
plt.xlabel('Energy Consumption')

plt.show()
```



3. Pre-process the data

As the graphs result shown above, we can find there are some outliers in appliances.
So I will deal with it below.

According to Turkey and John (1977), for "far out" outliers, k is equal to 3. So I'm going to base my data on this.

$Q1 = 50$, $Q3 = 100 \Rightarrow IQR = 100 - 50 = 50$

So, upper whisker = $100 + 3 \times 50 = 250$ and the lower whisker = $-3 \times 50 + 100 = -100$

As the result, we will remove the data which is higher than 250

```
In [ ]: outlier_rows = df[df['Appliances'] > 250]
        print(outlier_rows)
```

		date	Appliances	lights	T1	RH_1	\
11	2016-01-11	18:50:00	580	60	20.066667	46.396667	
12	2016-01-11	19:00:00	430	50	20.133333	48.000000	
31	2016-01-11	22:10:00	400	20	21.533333	44.966667	
32	2016-01-11	22:20:00	400	20	21.600000	44.766667	
33	2016-01-11	22:30:00	390	30	21.600000	44.560000	
...		
19704	2016-05-27	13:00:00	370	0	24.890000	47.730000	
19705	2016-05-27	13:10:00	280	0	25.033333	48.363333	
19732	2016-05-27	17:40:00	270	10	25.500000	46.596667	
19733	2016-05-27	17:50:00	420	10	25.500000	46.990000	
19734	2016-05-27	18:00:00	430	10	25.500000	46.600000	

		T2	RH_2	T3	RH_3	T4	...	T_out
\								
11		19.426667	44.400000	19.790000	44.826667	19.000000	...	5.98
12		19.566667	44.400000	19.890000	44.900000	19.000000	...	6.00
31		20.790000	44.360000	20.426667	45.933333	19.600000	...	5.65
32		20.890000	44.223333	20.500000	45.933333	19.696667	...	5.70
33		20.963333	43.963333	20.500000	45.790000	20.096667	...	5.75
...	
19704		26.500000	40.460000	28.730000	42.156667	24.500000	...	21.10
19705		26.528571	40.595714	28.496667	41.900000	24.500000	...	21.30
19732		25.628571	42.768571	27.050000	41.690000	24.700000	...	22.50
19733		25.414000	43.036000	26.890000	41.290000	24.700000	...	22.30
19734		25.264286	42.971429	26.823333	41.156667	24.700000	...	22.20

		Press_mm_hg	RH_out	Windspeed	Visibility	Tdewpoint	rv1
\							
11		734.433333	91.166667	5.833333	40.000000	4.62	8.827838
12		734.500000	91.000000	6.000000	40.000000	4.60	34.351142
31		735.883333	87.833333	6.166667	40.000000	3.72	29.978291
32		735.966667	87.666667	6.333333	40.000000	3.73	24.677065
33		736.050000	87.500000	6.500000	40.000000	3.75	9.310880
...	
19704		756.100000	60.000000	1.000000	24.000000	13.00	48.686116
19705		756.050000	59.833333	1.166667	23.666667	13.10	32.420348
19732		755.200000	56.333333	3.666667	25.333333	13.30	29.199117
19733		755.200000	56.666667	3.833333	26.166667	13.20	6.322784
19734		755.200000	57.000000	4.000000	27.000000	13.20	34.118851

	rv2	NSM	weekend
11	8.827838	67800.0	0
12	34.351142	68400.0	0
31	29.978291	79800.0	0
32	24.677065	80400.0	0
33	9.310880	81000.0	0
...
19704	48.686116	46800.0	0
19705	32.420348	47400.0	0
19732	29.199117	63600.0	0
19733	6.322784	64200.0	0
19734	34.118851	64800.0	0

[1531 rows x 31 columns]

```
In [ ]: # delete the outliers
df = df.drop(outlier_rows.index)

print(df['Appliances'].describe())
```

```
count    18204.000000
mean      72.153922
std       38.871238
min       10.000000
25%       50.000000
50%       60.000000
75%       90.000000
max       250.000000
Name: Appliances, dtype: float64
```

Now we can see the outliers are successfully removed.

Next, we will check the if there are any missing values in the data

```
In [ ]: missing_values = df.isna().sum()

print(missing_values)
```

```
date            0
Appliances      0
lights          0
T1              0
RH_1            0
T2              0
RH_2            0
T3              0
RH_3            0
T4              0
RH_4            0
T5              0
RH_5            0
T6              0
RH_6            0
T7              0
RH_7            0
T8              0
RH_8            0
T9              0
RH_9            0
T_out           0
Press_mm_hg     0
RH_out          0
Windspeed       0
Visibility       0
Tdewpoint       0
rv1             0
rv2             0
NSM             0
weekend         0
dtype: int64
```

As we can see, we don't have any missing value on the dataset

```
In [ ]: df.describe()
```

Out []:

	date	Appliances	lights	T1	RH
count	18204	18204.000000	18204.000000	18204.000000	18204.0000
mean	2016-03-20 16:38:07.343441152	72.153922	3.527247	21.687836	40.1776
min	2016-01-11 17:00:00	10.000000	0.000000	16.790000	27.0233
25%	2016-02-15 15:47:30	50.000000	0.000000	20.760000	37.2900
50%	2016-03-20 22:15:00	60.000000	0.000000	21.600000	39.5666
75%	2016-04-24 03:02:30	90.000000	0.000000	22.600000	42.9333
max	2016-05-27 17:30:00	250.000000	70.000000	26.200000	63.3600
std	NaN	38.871238	7.659798	1.606706	3.9465

8 rows × 31 columns

We can see that data anomalies in appliances have been removed. Below I will do the feature scaling.

Since 'date' has already been converted into two features : NSM and weekend. so we will exclude 'date' from dataset

```
In [ ]: df.drop('date', axis=1, inplace=True)
```

```
In [ ]: # using standard scaler to scale the data
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

# scale the dataset
df_scaled = scaler.fit_transform(df)

# Converts scaled data to a DataFrame and preserves the original column names
df_scaled = pd.DataFrame(df_scaled, columns=df.columns)
```

```
In [ ]: df_scaled.describe()
```

Out []:

	Appliances	lights	T1	RH_1	T2
count	1.820400e+04	1.820400e+04	1.820400e+04	1.820400e+04	1.820400e+04
mean	1.592515e-16	4.371610e-17	-4.996126e-17	1.080412e-15	-1.748644e-16
std	1.000027e+00	1.000027e+00	1.000027e+00	1.000027e+00	1.000027e+00
min	-1.599013e+00	-4.605009e-01	-3.048456e+00	-3.333197e+00	-1.927998e+00
25%	-5.699466e-01	-4.605009e-01	-5.774935e-01	-7.316996e-01	-6.950493e-01
50%	-3.126799e-01	-4.605009e-01	-5.467018e-02	-1.548090e-01	-1.572577e-01
75%	4.591201e-01	-4.605009e-01	5.677385e-01	6.982795e-01	5.470660e-01
max	4.575387e+00	8.678373e+00	2.808410e+00	5.874246e+00	4.377304e+00

8 rows × 30 columns

The dataset has successfully been scaled.

Then I will be splitting the data and divide it into two parts: the test set and the training set

In []:

```
from sklearn.model_selection import train_test_split

X = df_scaled.drop('Appliances', axis = 1)
y = df_scaled['Appliances']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

4. Implement, train and test prediction models

According to a study, I will use linear regression model as a prediction because it is considered relatively easy to implement, relatively less computationally intensive, high predictive power and high data availability (Fumo & Rafe 2015, p.333).

In []:

```
# linear regression
from sklearn.linear_model import LinearRegression
linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)
y_pred_linear = linear_reg.predict(X_test)

from sklearn.metrics import mean_squared_error, mean_absolute_error

mse_linear = mean_squared_error(y_test, y_pred_linear)
rmse_linear = np.sqrt(mse_linear)
mae_linear = mean_absolute_error(y_test, y_pred_linear)
linear_trainScore = linear_reg.score(X_train, y_train)

print("Linear Regression Model:")
print("Mean Squared Error (MSE):", mse_linear)
```

```
print("Root Mean Squared Error (RMSE):", rmse_linear)
print("Mean Absolute Error (MAE):", mae_linear)
print("Train Score of Linear Regression:", linear_trainScore)
```

Linear Regression Model:

Mean Squared Error (MSE): 0.7389807430157602

Root Mean Squared Error (RMSE): 0.859639891475355

Mean Absolute Error (MAE): 0.5610255129608029

Train Score of Linear Regression: 0.26859173550003856

According to the study, Random Forest outperforms other widely used classifiers such as Artificial Neural Networks and Support Vector Machines in energy consumption prediction (Chen et al. 2019, p. 957). So we also use random forest model to predict the energy consumption

```
In [ ]: from sklearn.ensemble import RandomForestRegressor

random_forest = RandomForestRegressor(n_estimators=100, random_state=42)
random_forest.fit(X_train, y_train)
y_pred_rf = random_forest.predict(X_test)
mse_rf = mean_squared_error(y_test, y_pred_rf)
rmse_rf = np.sqrt(mse_rf)
mae_rf = mean_absolute_error(y_test, y_pred_rf)
rf_trainScore = random_forest.score(X_train, y_train)

print("Random Forest Model:")
print("Mean Squared Error (MSE):", mse_rf)
print("Root Mean Squared Error (RMSE):", rmse_rf)
print("Mean Absolute Error (MAE):", mae_rf)
print("Train Score of Random Forest Regression:", rf_trainScore)
```

Random Forest Model:

Mean Squared Error (MSE): 0.40602471396744605

Root Mean Squared Error (RMSE): 0.637200685787018

Mean Absolute Error (MAE): 0.3702047117140477

Train Score of Random Forest Regression: 0.9453180787804506

V E et al.(2021) used a KNN model for prediction and they concluded that the efficiency of KNN is optimal for practical applications. So I use their methodology here to predict energy consumption

```
In [ ]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error

knn_reg = KNeighborsRegressor(n_neighbors=5)
knn_reg.fit(X_train, y_train)

y_pred_knn = knn_reg.predict(X_test)

mse_knn = mean_squared_error(y_test, y_pred_knn)
rmse_knn = np.sqrt(mse_knn)
mae_knn = mean_absolute_error(y_test, y_pred_knn)
knn_train_score = knn_reg.score(X_train, y_train)

print("KNN Regression Model:")
print("Mean Squared Error (MSE):", mse_knn)
print("Root Mean Squared Error (RMSE):", rmse_knn)
```

```
print("Mean Absolute Error (MAE):", mae_knn)
print("Train Score of KNN Regression:", knn_train_score)
```

KNN Regression Model:

Mean Squared Error (MSE): 0.5633286438394922

Root Mean Squared Error (RMSE): 0.750552225924014

Mean Absolute Error (MAE): 0.4508702888796436

Train Score of KNN Regression: 0.63918955704182

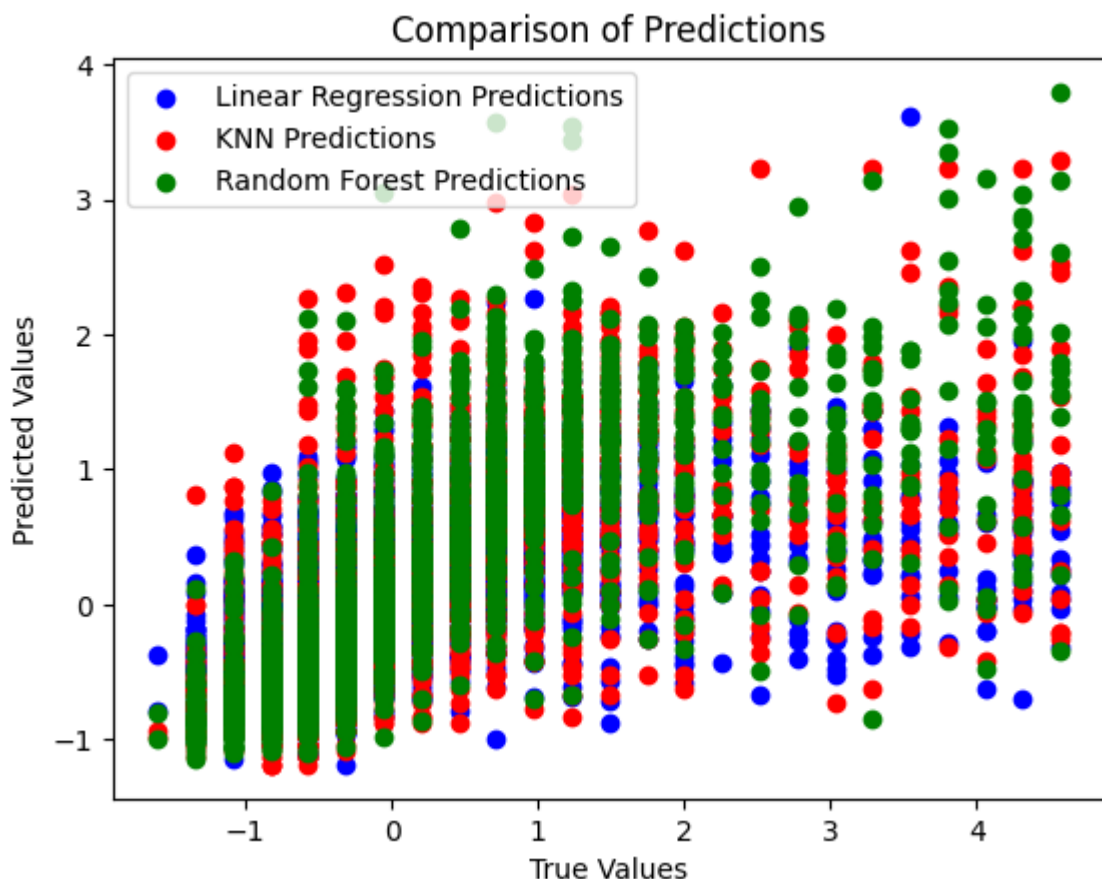
5. Compare the results from all candidate models, choose the best model, justify your choice and discuss the results

MSE, MAE and RMSE are convenient measures of "mean error" used to assess the degree of variability in the data (Lyu et al. 2022, p.1488). So we visualise the data based on the above results for comparison

```
In [ ]: plt.scatter(y_test, y_pred_linear, color='blue', label='Linear Regression')
plt.scatter(y_test, y_pred_knn, color='red', label='KNN Predictions')
plt.scatter(y_test, y_pred_rf, color='green', label='Random Forest Predictions')

plt.title("Comparison of Predictions")
plt.xlabel("True Values")
plt.ylabel("Predicted Values")

plt.legend()
plt.show()
```



```
In [ ]: model_labels = ['Linear Regression', 'Random Forest', 'KNN']

mse_scores = [0.7389807430157602, 0.40602471396744605, 0.5633286438394922]
rmse_scores = [0.859639891475355, 0.637200685787018, 0.750552225924014]
```

```

mae_scores = [0.5610255129608029, 0.3702047117140477, 0.4508702888796436]
r2_scores = [0.26859173550003856, 0.9453180787804506, 0.63918955704182]

plt.figure(figsize=(12, 8))

# MSE
plt.subplot(2, 2, 1)
plt.bar(model_labels, mse_scores, color=['skyblue', 'lightgreen', 'salmon'])
plt.title('Mean Squared Error (MSE)')
plt.ylabel('MSE')

# RMSE
plt.subplot(2, 2, 2)
plt.bar(model_labels, rmse_scores, color=['skyblue', 'lightgreen', 'salmon'])
plt.title('Root Mean Squared Error (RMSE)')
plt.ylabel('RMSE')


# MAE
plt.subplot(2, 2, 3)
plt.bar(model_labels, mae_scores, color=['skyblue', 'lightgreen', 'salmon'])
plt.title('Mean Absolute Error (MAE)')
plt.ylabel('MAE')

# R2
plt.subplot(2, 2, 4)
plt.bar(model_labels, r2_scores, color=['skyblue', 'lightgreen', 'salmon'])
plt.title('R2 Score')
plt.ylabel('R2')

plt.tight_layout()

plt.show()

```



We can conclude that random forest has the smallest model prediction error and therefore the model is better at predicting. As the result, random forest model works

the best in predicting the energy consumption. In terms of R^2 score, random forest reached 0.8

6. Reflect on what you have learned by completing this assignment and how to improve the models

By completing this assignment I have learnt how to visualise data, pre-process data and use implement machine learning models such as linear regression, random forests and k-nearest neighbours. In fact, I took a similar class last year called Computer Vision, which had roughly the same steps except for the data visualisation. In terms of model improvement and optimisation, I think I have room for improvement. For example, I can normalise the data. For model selection, I can also try more models to get the best model. In terms of dataset, I can try to add new attributes that will help the models to capture patterns and trends in the data more accurately.

7. References

- [1] Candanedo, LM, Feldheim, V & Deramaix, D 2017, 'Data driven prediction models of energy use of appliances in a low-energy house', *Energy and Buildings*, vol. 140, pp. 81–97.
- [2] Chen, YT, Piedad, E & Kuo, CC 2019, 'Energy consumption load forecasting using a level-based random forest classifier', *Symmetry (Basel)*, vol. 11, no. 8, pp. 956–965.
- [3] Fumo, N & Rafe Biswas, MA 2015, 'Regression analysis for prediction of residential energy consumption', *Renewable & Sustainable Energy Reviews*, vol. 47, pp. 332–343.
- [4] Lyu, Z, Yu, Y, Samali, B, Rashidi, M, Mohammadi, M, Nguyen, TN & Nguyen, A 2022, 'Back-Propagation Neural Network Optimized by K-Fold Cross-Validation for Prediction of Torsional Strength of Reinforced Concrete Beam', *Materials*, vol. 15, no. 4, pp. 1477–1498.
- [5] Tukey, JW & John W, 1977, *Exploratory data analysis*, Addison-Wesley Pub. Co., U.S.A.
- [6] V E, S, Shin, C & Cho, Y 2021, 'Efficient energy consumption prediction model for a data analytic-enabled industry building in a smart city', *Building Research and Information: The International Journal of Research, Development and Demonstration*, vol. 49, no. 1, pp. 127–143.